

# NoC-based Multi-FPGA Application Mapping: A Case Study Accelerating Krylov Sequence Computations over GF(2)

*Abstract—*

## I. INTRODUCTION

### A. Relevant focus areas for FSP

- Design automation for multi-FPGA and heterogeneous systems.
- Overlays: FPGA-aware CONNECT[1] + NoC + Processors?
- Mapping approaches and tools for heterogeneous FPGAs.
- HLS. High-level compilation and languages, design automation tools that raise the abstraction level when designing for (heterogeneous) FPGAs and reconfigurable systems and standardized target platforms.
- Application case studies.

### B. Organization

## II. RELATED WORK

### A. Network-on-Chip

Network-on-Chip has been a topic of immense interest in the last decade. A rich body of work exists on several themes concerning NoC, such as custom topology, routing, micro-architecture, power, placement, verification among several more. Multi-FPGA partitioning for prototyping of large designs dates back to the pre-NoC era, when the FPGA was a lot more resource constrained. Ouass et. al [2] developed "SPARCS", a synthesis and partitioning tool for multi-FPGA systems, with shared memory or direct channel communication between partitioned tasks. Roy-Neogi et. al [3] suggested a genetic algorithm based partitioning method for multiple FPGAs within resource and timing constraints. Papamichael et. al [1] developed an automated NoC generator tool tailor-made for FPGAs. Liu et. al [4] have developed a multi-FPGA emulation framework, with MDT transceivers for communication. However, the system only "mimics" NoC, and does not implement it on FPGA. We are not aware of any prior work which extends the NoC methodology itself in the multi-FPGA scenario, similar to our work.

### B. Boolean Matrix Vector Multiplication

Boolean Matrix Vector Multiplication has also been a topic of interest, especially with its application in Number Field Sieve (NFS) for integer factorization [5]. Bernstein [6] proposed a circuit-based implementation of the matrix-step in the Number Field Sieve. Geiselmann et. al [7] proposed a hardware implementation of the "mesh routing" algorithm using several interconnected ASICs. Bajracharya et. al [8] proposed an implementation of the above mesh routing circuit on a reconfigurable architecture. In mesh routing, a matrix  $A$  of size  $n \times n$ , requires a mesh of size  $m \times m$  nodes, where  $m = \sqrt{n \cdot h}$ ,  $h$  being the maximum number of non-zero bits in any column of matrix  $A$ . The  $m \times m$  nodes are further divided into  $D$  blocks, each of size  $h \times h$ , which are initialized with the input matrix  $v$  and would produce the product vector  $v'$  at the end of the algorithm. The  $h$  nodes in a block  $i$  are required to store the

addresses (requiring  $\log(n)$  bits) of the non-zero bits in column  $i$ . For computing the product, each non-zero block would "route" a message to target blocks corresponding to the  $h$  addresses of that block, with each receiving block  $i$ , flipping its single-bit product value (initialized to 0 before routing starts) corresponding to  $i$ -th bit in  $v'$  on reception a message.

Our work offers a different perspective to the BMVM in Matrix Step, which has largely been focussed around the mesh routing, and differs from it in several ways. First, mesh routing has been designed to handle extremely sparse matrices, but our method is better suited for handling dense matrices. This is because, for dense matrices, mesh routing is required to store  $O(n^2)$  and preform  $O(n^2)$  operations. In contrast, our method requires  $n/k$  or  $O(n)$  processing elements, and  $O(n^2/k^2)$   $k$ -bit operations. Our pre-processing stage significantly helps in reducing the number of operations per iteration of BMVM, making it more suitable to handle iterative BMVM in case of dense matrices. We are yet to modify our algorithm for the case of "sparse" matrices. Second, each message in mesh routing corresponds to a single bit operation, while a message in our technique correspond to  $k$ -bit operations. Third, mesh routing has been specifically devised for mesh-based networks, with a tailor-made routing strategy. Our technique, on the other hand, is applicable to a flexible NoC-based interconnection, and can be used with low-radix, as well as high-radix topologies. As our results observe, the chosen topology can have a significant bearing on the overall performance, and this flexibility allows the designer to obtain a fine balance between cost and performance.

There has also been a rich history of literature on accelerating sparse matrix vector multiplication using a multitude of techniques, such as graph partitioning etc, but from the purview of parallel multi-processors and not custom hardware [9], [10].

## III. BOOLEAN MATRIX VECTOR MULTIPLICATION

Matrix-vector multiplication is central to a large number of scientific computing applications. Boolean matrix vector multiplication is defined over the smallest finite field GF(2). Boolean matrix vector product has important applications in coding theory [11], cryptanalysis [5] and image compression [12], among many others. Several secure systems today rely on the computational intractability of factoring large numbers [13]. Lenstra et. al [5], proposed Number Field Sieve (NFS), perhaps the fastest known algorithm till date for factoring large integers. NFS, on practical applications, would still require several months of computational time to be completely solved on several computing clusters. There are four major steps in the NFS algorithm:

- 1) Polynomial selection
- 2) Sieving
- 3) Matrix Step
- 4) Square root

*Sieving* and *Matrix Step* are considered to be the most time-consuming steps in the NFS [14]. In particular, the Matrix Step involves finding linear dependencies in a large boolean matrix. Block Wiedemann [15] algorithm is often used for this purpose. This involves repeated multiplying of a very large boolean matrix  $A$  (matrix size in the range  $10^6 - 10^{11}$ ), starting with a random boolean vector  $v_i$  resulting in product vectors  $A.v_i, A^2.v_i, \dots, A^k.v_i$ . This is done for a large number of random boolean vectors  $v_i$ , with  $i$  ranging from 1 to  $k$ , where  $k = 2D/K$ . Here,  $D$  is the column size of matrix  $A$  and  $K$  is known as the “blocking factor”.

In general, the boolean matrix-vector multiplication (BMVM) is considered to be quadratic w.r.t. to the size (column/row size) of the matrix  $A$ . However, since we are repeatedly using the matrix  $A$  for computing products with random vectors over thousands of iterations, any improvement in the computation time for computing product by exploiting the structure of matrix  $A$ , would directly reflect in the runtime of the *Block Wiedemann* algorithm, and therefore, also of *NFS*. Geiselmann et. al [7] have made certain assumptions on the *sparsity* of matrix  $A$ , and exploited it for accelerating the boolean matrix-vector product using custom hardware.

In this work, we do not wish to make any assumptions on the sparseness of matrix  $A$  as in [7], since we want our technique to be more generally applicable to other applications. We do, however, believe that exploiting the sparsity structure of matrix would be necessary for scalability of our technique, and briefly discuss it in section II. We have left these ideas as a part of future work. Our technique is based on the recently proposed combinatorial approach for matrix vector multiplication by Ryan Williams [16]. This approach allows matrix-vector multiplication to be performed sub-quadratically in  $O((n/k)^2)$  on a  $\log(n)$ -word RAM, with a single pre-processing step of  $O(n^{2+\epsilon})$ , when  $k = \epsilon \log(n)$ .

We describe this technique in section III-A (for more details on the algorithm, refer to [16]), and then report our hardware implementation details in section IV.

#### A. Sub-quadratic Method to Matrix Vector Multiplication

We now describe the algorithm for Boolean Matrix Vector Multiplication, which allows for sub-quadratic multiplication, with some pre-processing. The algorithm is based on [16], and has two phases: (i) pre-processing and (ii) computation phase. The pre-processing step uses the given input  $n \times n$  matrix  $A$  to build look-up tables. The computation step uses the generated look-up tables to perform the actual multiplication with the given input vector  $v$ . We now individually describe the steps involved in the two phases. Note that we use a modified convention over [16] for simplified understanding.

##### 1) Pre-processing Phase:

- 1) Partition the input matrix  $A$  into blocks (tiles) of size  $k \times k$  bits as shown in Figure 1, for a given user defined parameter  $k$ . The value of this parameter  $k$  has important ramifications on the requirements of the number of look-up tables and computing nodes, size of the look-up tables and the number of operations during pre-processing and computation phases. These relations would become clear by the end of this section.
- 2) Construct  $n/k$  loop-up tables (LUTs), labelled as  $LUT_1, LUT_2, \dots, LUT_{n/k}$ , each of which further consists

of  $2^k$  partitions. These partitions are indexed from 0 to  $2^k - 1$ . Each partition consists of  $n/k$  words of size  $k$  bits.

- 3) Populate the LUTs such that the  $p$ -th partition of  $LUT_i$  consists of  $k$ -bit words corresponding to values  $A_{1,i}.b_p, A_{2,i}.b_p, \dots, A_{n/k,i}.b_p$ . Here  $b_p$  is a  $k$ -bit vector representation corresponding to the decimal value  $p$ . This means that  $b_0$  is an *all-zero* vector and  $b_{2^k-1}$  is an *all-one* vector. The structure of the  $LUT_i$  is also shown in Figure 1.

Note that pre-processing step is equivalent to pre-computing and storing all possible products of the blocks of matrix  $A$ , i.e. for  $A_{1,1}, A_{1,2}, \dots, A_{n/k,n/k}$ . Since there are  $2^k$  such possible products for each block, and  $n^2/k^2$  blocks, a total of  $2^k \cdot (n^2/k^2)$  matrix products are required during pre-processing. Since each of this product requires  $O(k^2)$  computational operations, and hence, the total pre-processing can be achieved in  $O(n^2 \cdot 2^k)$  operations. Also if the vector  $v$  and the product vector  $v' = A.v$  are also partitioned into  $n/k$  sub-vectors,

$$\text{each of size } k\text{-bits, as } v = \begin{pmatrix} v_1 \\ v_2 \\ \vdots \\ v_{n/k} \end{pmatrix} \text{ and } v' = \begin{pmatrix} v'_1 \\ v'_2 \\ \vdots \\ v'_{n/k} \end{pmatrix}, \text{ then the}$$

sub-vector  $v'_i$  can also be expressed as:

$$v'_i = (A_{i,1}.v_1) \vee (A_{i,2}.v_2) \vee \dots (A_{i,n/k}.v_{n/k}) \quad (1)$$

In equation 1, the operator  $\vee$  performs bit-wise *xor*-operation on the  $k$ -bit partial products. This is the essence of the computing phase described next.

##### 2) Computing Phase:

- 1) Partition the input vector  $v$  into  $n/k$   $k$ -bit sub-vectors

$$v_1, v_2, \dots, v_{n/k}, \text{ s.t. } v = \begin{pmatrix} v_1 \\ v_2 \\ \vdots \\ v_{n/k} \end{pmatrix}. \text{ Let } v' \text{ be another } n\text{-bit vector.}$$

Partition  $v'$  into  $k$ -bit sub-vectors in a similar manner, as  $v' =$

$$\begin{pmatrix} v'_1 \\ v'_2 \\ \vdots \\ v'_{n/k} \end{pmatrix}, \text{ and initialize it to 0.}$$

- 2) Look-up the partition  $p_i$  in the look-up table  $LUT_i$  (obtained from pre-processing) corresponding to the sub-vector  $v_i$  for all  $i$  in  $\{1, 2, \dots, n/k\}$ .
- 3) For the  $j$ -th ( $j$  ranging from 1 to  $n/k$ ) word  $e_j$  in partition  $p_i$ , update  $v'_j$  as bit-wise *xor* of current  $v'_j$  with  $e_j$ . Do this for all partitions  $p_i$ , obtained in step 2.

$$4) \text{ The product } A.v \text{ is then given by } v' = \begin{pmatrix} v'_1 \\ v'_2 \\ \vdots \\ v'_{n/k} \end{pmatrix}.$$

It can be inferred easily that the computing phase requires  $n^2/k^2$  bit-wise *xor*-operations of size  $k$ -bits. The matrix-vector multiplication can, therefore, be computed in  $O(n^2/k^2)$ . In [16], the authors have used  $k = \epsilon \log_2(n)$ , with  $\epsilon < 0.5$ , so that the multiplication can be carried out in  $O(n^2/(\epsilon \log(n))^2)$  on a  $\log(n)$ -word RAM.

We summarize the important analysis related to the algorithm, when used with parameter  $k$  here:

Word width of LUT =  $k$

Number of LUTs =  $n/k$

Size of each LUT (in words) =  $2^k \cdot (n/k)$

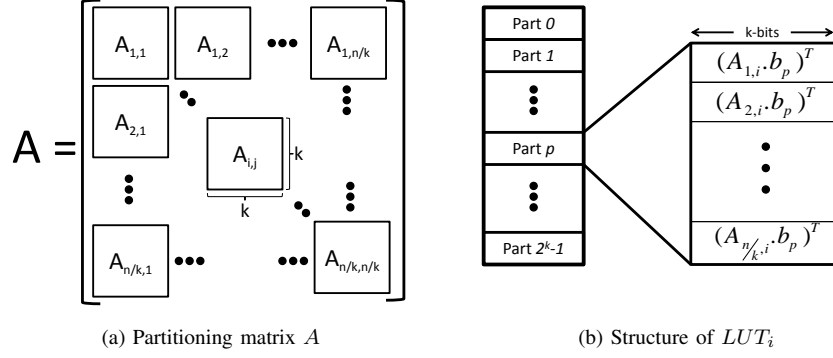


Fig. 1: Preprocessing of Matrix  $A$

Number of word operations required =  $n^2/k^2$

It is observable that increasing the value of  $k$  would reduce the computation time due to the reduced number of operations, but would, in turn, also increase the memory requirement (in LUTs) and the pre-processing time (occurring in  $O(2^k \cdot (n^2/k^2))$ ) exponentially. A careful selection of the parameter  $k$  is thus required for achieving the right balance between the three. Also note that the extreme case when  $k = n$  requires a single LUT storing  $n$ -bit product vectors corresponding to all possible  $2^n$  combinations of input vector. Any multiplication would only require a single look-up into the corresponding entry. However, this is exponential in memory and pre-processing time and cannot be applied to matrices exceeding a few 10s of bits in size.

### B. Implementation Details

In this section, we shall describe the implementation details of the algorithm described in Section III-A. The goal of our implementation is to have a *flexible* and *scalable* software library framework for FPGA-accelerated iterative boolean matrix multiplication using the above algorithm.

1) *Pre-processing*: As described in section III-A, the algorithm is composed of two phases: the pre-processing phase and the computation phase. Since the pre-processing phase is required only once for a given input matrix, which is assumed to be re-used multiple times over several iterations in the target applications, we apply this phase completely on the software end using MATLAB. The entire source code, along with usage details, for this step is available in section ???. The script takes the matrix size  $n$ ,  $\epsilon$  (s.t.  $k = \epsilon \cdot \log(n)$ ) and folding factor as an input parameter. We describe the concept of "folding", which is used by the folding factor, later in this section.

It must be noted that although the authors in [16] use  $\epsilon < 0.5$  for reasons mentioned already, we make no such assumption and leave it to the user to best determine the appropriate value of  $\epsilon$ . The code produces several '.txt' and '.data' files as output, corresponding to the look-up tables used in the FPGA and software implementation of the algorithm, respectively and having the same structure as described in Figure 1(b). The  $LUTs$  are generated by multiplying each blocks (tiles) of size  $k \times k$  of the input matrix  $A$  by all possible vector combinations of size  $k$ , and storing the output values. In the current code, the matrix  $A$ , of size  $n \times n$ , is randomly generated with




Fig. 2: FPGA internal structure.

uniform probability of 1's and 0's. The matrix size is also assumed to be a power of 2, which can be easily modified to make it more general.

2) *Hardware implementation*: We now start discussing the FPGA implementation of the algorithm. As evident in section III-A, the algorithm is memory-intensive, rather than computation-intensive, due to the large sizes of the look-up-tables. To preserve data locality and improve computation time, we would like the  $LUTs$  to completely reside on the FPGA hardware. In order to best utilize the available memory resources in the FPGA, we first observe that modern FPGAs are provisioned with abundant Block RAM (BRAM) resources. For instance, *Xilinx Vitex 6* (used in this work) has a large number of 36Kb Block RAMs, totalling upto 38Mb of internal BRAM storage. The internal structure of a typical FPGA is shown in Figure 6. It can be observed that the Block RAMs are broadly distributed over the entire FPGA fabric.

As mentioned earlier, the memory requirement of a single  $LUT$  is  $(n/k) \cdot 2^k$  words (with each word storing  $k$  bits). We would ideally like the  $LUTs$  to be implemented automatically using Block RAMs.

It must be noted that the look-up tables, implemented in the form of Block RAMs are required to "exchange information" based on the algorithm described in Section III-A. The BRAMs could be situated far off from each other on an FPGA, or even reside

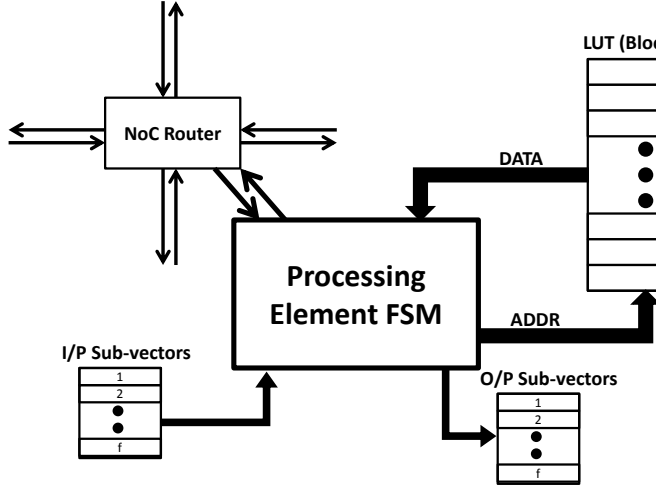


Fig. 3: Processing element.

on different FPGAs altogether, in a multi-FPGA implementation (memory constraint of FPGA limits the size and number of LUTs that can reside on a single FPGA). Network-on-chip would not only help in seamless communication between the LUTs, but would also help improve the clock speed. In our FPGA implementation, with each of the  $n/k$  LUTs obtained from pre-processing, we associate a “computing node”. These nodes can communicate with each other using  $k$ -bit messages over an NoC and these nodes are indexed from 1 to  $n/k$ . A compute node  $i$  also maintains a  $k$ -bit register  $v'_i$ , initialized to 0. A compute node indexed  $i$  would look-up the partition  $p_i$  in the look-up table  $LUT_i$  corresponding to the sub-vector  $v_i$ . The compute node  $i$  sends  $(n/k)$   $k$ -bit values in the partition  $p_i$  to the corresponding compute nodes. Every compute node also would receive  $n/k$  such  $k$ -bit messages. The register  $v'_i$  corresponding to node  $i$  is computed by bit-wise *xor* of all the messages received by node  $i$ . The product  $A.v$  is then given by  $v'$

$$= \begin{pmatrix} v'_1 \\ v'_2 \\ \vdots \\ v'_{n/k} \end{pmatrix}.$$

We use a NoC generated by CONNECT[1] for our implementation in which the computing nodes are connected as the “processing elements” (PEs) at the input/output ports of the generated NoC. These nodes use the interface provided by CONNECT to exchange messages. It is important to ensure that while multiple such messages may simultaneously attempt to update a particular product sub-vector  $v'_i$ , the updates are appropriately serialized to maintain correctness. Since only one flit can be injected and ejected in a single cycle in the NoC, this constraint is automatically ensured. our implementation supports the following “Network and Router Options” for NoC generated using CONNECT (topology and number of endpoints is user-specified):

- **Router Type:** Simple Input Queued (IQ)
- **Flow Control Type:** Peek Flow Control
- **Flit Data Width:** 16
- **Flit Buffer Depth:** 8
- **Allocator:** Separable Input first Round-Robin

Since the number of computing nodes required grows linearly

with the matrix size (as  $n/k$ ), the number of ports required for the NoC would also grow linearly. This could pose an exorbitant requirement on the number of NoC ports since most topologies do not scale well to a very large number of endpoints. To this end, we implement “folding” of computing nodes, which allows multiple processing elements connected to a single endpoint of an NoC to serve as multiple computing nodes. The extent of such folding, referred to as the “folding factor”  $f$ , is a user-specified parameter in our implementation. Our folding follows a cyclic assignment, i.e. an instance with  $n$  computing nodes and  $n/f$  processing elements or endpoints, would assign computing nodes  $1, f+1, 2f+1, \dots, (n-f)+1$  computing nodes to the processing element 1. Since a single processing element is now required to send and receive  $f^2$  times the number of messages to the case where there is no folding with same number of processing elements (or ports in the network), the message traffic would also increase, resulting in higher network utilization. Since only one message can be sent/received on a single NoC port, the folded design would exploit lower concurrency, thereby lowering performance. The folding parameter is also specified while generating the LUTs in the pre-processing step, which results in large memory sizes compared to the case of no folding. The caveat of large folding, however, is that large LUT could also lower the maximum operational frequency of the design, thereby further lowering its overall performance. These effects are discussed in greater detail in Chapter ?? . Figure 7 shows a single processing element (multiplication node) with folding factor  $f$  in our implementation.

Since a number of our target applications require computation of several iterations of boolean vector multiplications, viz.  $A.v, A^2.v, \dots, A^k.v$ , given an input vector  $v$ , we provision our processing elements to support iterative matrix vector product. The number of such iterations is again provided as parameter by the user. Support for iterative multiplication would require the processing elements to perform multiplication of matrix  $A$  with input vector  $v'_k$  in the  $k$ th iteration, where  $v'_k$  is the resulting product vector after the  $(k-1)$ th iteration. Since the resulting sub-vectors  $v'_k$  are local to the processing elements in the  $k$ th iteration, this task seems trivial. However, iterative multiplication poses a unique challenge of *synchronization* of processing nodes after every iteration. This means that a processing element should only commence the  $k$ th iteration after ensuring that all other processing elements are also ready for starting  $k$ th iteration. If this is not the case and a processing element (which has already sent and received all messages of previous iteration) jumps to next iteration before others, the messages sent by this PE could be misinterpreted by other PEs to belong to previous iteration, resulting in incorrect computation of product values. One solution could be to use additional bits in messages to indicate the *id* for the iteration to which the message belongs. This would allow PEs to proceed to next iteration without waiting for others but this is expensive as the flit width is limited and this would also require buffering of incoming messages of future iterations. Other option could also be to use special messages to broadcast information to all PEs whenever a PE has completed an iteration. Each PE would maintain a list of completed PEs and would proceed to next iteration only after ensuring all others are also ready to proceed. While this seems more feasible, it also has performance penalty associated with the broadcast messages after every iteration.

In our approach, we observe that the behaviour of the network




Fig. 4: Top module for Boolean Matrix Vector Multiplication (BMVM).

after in each iteration would be exactly identical if the PEs are exactly synchronized. Therefore, if we know the maximum number of cycles required for a PE to complete an iteration, all PEs could wait to synchronize at this cycle count after every iteration, thereby eliminating the need for additional broadcast messages. It is important to note that this count is dependent on the topology and routing policy used. We specify the count as a parameter to the processing element in our design, which could be determined by the user by simulating one iteration of the multiplication on the software. Our processing elements display the cycle counts at which the processing elements are done with one iteration (when they have sent and received all messages pertaining to that iteration), and the maximum of this count can be specified in that parameter. The source code of our processing element is included in ?? . Other parameters used by this module are specified and explained in ?? . Figure 8 shows the top level user module along with their module names in our design, consisting of several instances of processing elements connected by a CONNECT generated NoC. It must be noted that the NoC can itself be partitioned using our script, in which case the top module would consist of multiple instances for NoC corresponding to different partitions, with serial wire and flow control signals between them. We have only tested our partitioned NoC for boolean vector multiplication on a single FPGA, with wires corresponding to different partitions connected internally. Since the partitioned NoC itself has been tested successfully on separate Virtex 5 boards, we believe that multiplier design should also work for multi-FPGA case. The interface to the top module is kept as simple as providing the input vector followed by a reset signal as input, and obtaining the output vector on the done signal at the output.

In our final step, we wish to encapsulate our multiplication hardware within a software library to make it usable for target applications, mostly running on software. For this purpose, we use Reusable Integration Framework for FPGA Accelerators (RIFFA 2.0) [17], which provides a simple open-source integration framework for FPGA-CPU communication over high-speed PCIe links. On the software end, RIFFA provides a simple API to open/close an FPGA connection, send/receive data from FPGA and to send reset the FPGA user module. Although we use a single

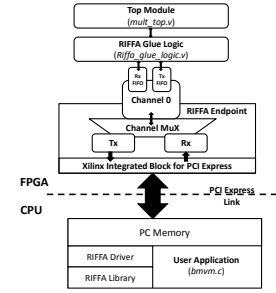


Fig. 5: BMVM framework with RIFFAa.

channel (corresponding to network sockets and parallel user cores), RIFFA can support upto 12 channels. The data transfers in RIFFA occur through PCIe transactions, with FPGA acting as the DMA bus master. For more details, the reader can refer to [17].

We demonstrate the usage using a simple C application shown in ?? . The application resets the FPGA and sends an input vector, on which the multiplication is performed by the FPGA, before waiting to receive the output product. Since the application is written for a 32-bit processor, appropriate precautions are taken to make it compatible with the *little endianess* of the RIFFA hardware modules supporting 64-bits. We also implement the necessary glue logic for the hardware interface between RIFFA endpoint and multiplier module as shown in Figure 9. The code for this module is available in ?? .

#### IV. IMPLEMENTATION DETAILS

In this section, we shall describe the implementation details of the algorithm described in Section III-A. The goal of our implementation is to have a *flexible* and *scalable* software library framework for FPGA-accelerated iterative boolean matrix multiplication using the above algorithm.

##### A. Pre-processing

As described in section III-A, the algorithm is composed of two phases: the pre-processing phase and the computation phase. Since the pre-processing phase is required only once for a given input matrix, which is assumed to be re-used multiple times over several iterations in the target applications, we apply this phase completely on the software end using MATLAB. The entire source code, along with usage details, for this step is available in section ?? . The script takes the matrix size  $n$ ,  $\epsilon$  (s.t.  $k = \epsilon \cdot \log(n)$ ) and folding factor as an input parameter. We describe the concept of "folding", which is used by the folding factor, later in this section.

It must be noted that although the authors in [16] use  $\epsilon < 0.5$  for reasons mentioned already, we make no such assumption and leave it to the user to best determine the appropriate value of  $\epsilon$ . The code produces several '.txt' and '.data' files as output, corresponding to the look-up tables used in the FPGA and software implementation of the algorithm, respectively and having the same structure as described in Figure 1(b). The *LUTs* are generated by multiplying each blocks (tiles) of size  $k \times k$  of the input matrix  $A$  by all possible vector combinations of size  $k$ , and storing the output values. In the current code, the matrix  $A$ , of size  $n \times n$ , is randomly generated with




Fig. 6: FPGA internal structure.

uniform probability of 1's and 0's. The matrix size is also assumed to be a power of 2, which can be easily modified to make it more general.

### B. Hardware implementation

We now start discussing the FPGA implementation of the algorithm. As evident in section III-A, the algorithm is memory-intensive, rather than computation-intensive, due to the large sizes of the look-up-tables. To preserve data locality and improve computation time, we would like the *LUTs* to completely reside on the FPGA hardware. In order to best utilize the available memory resources in the FPGA, we first observe that modern FPGAs are provisioned with abundant Block RAM (BRAM) resources. For instance, *Xilinx Vitex 6* (used in this work) has a large number of 36Kb Block RAMs, totalling upto 38Mb of internal BRAM storage. The internal structure of a typical FPGA is shown in Figure 6. It can be observed that the Block RAMs are broadly distributed over the entire FPGA fabric.

It must be noted that the loop-up tables, implemented in the form of Block RAMs are required to “exchange information” based on the algorithm described in Section III-A. The BRAMs could be situated far off from each other on an FPGA, or even reside on different FPGAs altogether, in a multi-FPGA implementation (memory constraint of FPGA limits the size and number of LUTs that can reside on a single FPGA). Network-on-chip would not only help in seamless communication between the LUTs, but would also help improve the clock speed. In our FPGA implementation, with each of the  $n/k$  *LUTs* obtained from pre-processing, we associate a “computing node”. These nodes can communicate with each other using  $k$ -bit messages over an NoC and these nodes are indexed from 1 to  $n/k$ . A compute node  $i$  also maintains a  $k$ -bit register  $v'_i$ , initialized to 0. A compute node indexed  $i$  would look-up the partition  $p_i$  in the look-up table  $LUT_i$  corresponding to the sub-vector  $v_i$ . The compute node  $i$  sends  $(n/k)$   $k$ -bit values in the partition  $p_i$  to the corresponding compute nodes. Every compute node also would receive  $n/k$  such  $k$ -bit messages. The register  $v'_i$  corresponding to node  $i$  is computed by bit-wise *xor* of all the messages received by node  $i$ . The product  $A.v$  is then given by  $v'$

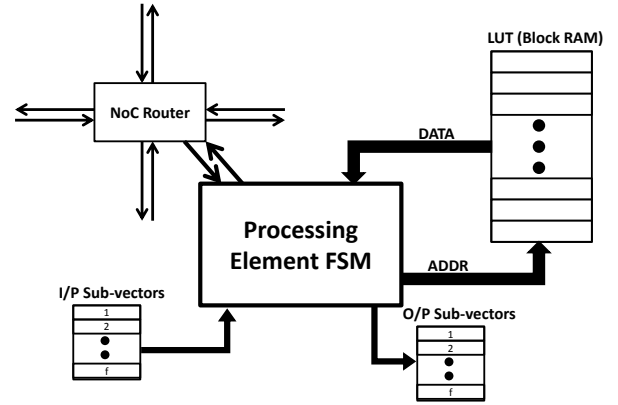


Fig. 7: Processing element.

$$= \begin{pmatrix} v'_1 \\ v'_2 \\ \vdots \\ v'_{n/k} \end{pmatrix}.$$

We use a NoC generated by CONNECT[1] for our implementation in which the computing nodes are connected as the “processing elements” (PEs) at the input/output ports of the generated NoC. These nodes use the interface provided by CONNECT to exchange messages. It is important to ensure that while multiple such messages may simultaneously attempt to update a particular product sub-vector  $v'_i$ , the updates are appropriately serialized to maintain correctness. Since only one flit can be injected and ejected in a single cycle in the NoC, this constraint is automatically ensured. Our implementation supports the following “Network and Router Options” for NoC generated using CONNECT (topology and number of endpoints is user-specified):

- **Router Type:** Simple Input Queued (IQ)
- **Flow Control Type:** Peek Flow Control
- **Flit Data Width:** 16
- **Flit Buffer Depth:** 8
- **Allocator:** Separable Input first Round-Robin

Since the number of computing nodes required grows linearly with the matrix size (as  $n/k$ ), the number of ports required for the NoC would also grow linearly. This could pose an exorbitant requirement on the number of NoC ports since most topologies do not scale well to a very large number of endpoints. To this end, we implement “folding” of computing nodes, which allows multiple processing elements connected to a single endpoint of an NoC to serve as multiple computing nodes. The extent of such folding, referred to as the “folding factor”  $f$ , is a user-specified parameter in our implementation. Our folding follows a cyclic assignment, i.e. an instance with  $n$  computing nodes and  $n/f$  processing elements or endpoints, would assign computing nodes  $1, f+1, 2f+1, \dots, (n-f)+1$  computing nodes to the processing element 1. Since a single processing element is now required to send and receive  $f^2$  times the number of messages to the case where there is no folding with same number of processing elements (or ports in the network), the message traffic would also increase, resulting in higher network utilization. Since only one message can be sent/received on a single NoC port, the folded design would exploit lower concurrency, thereby lowering performance. The

folding parameter is also specified while generating the LUTs in the pre-processing step, which results in large memory sizes compared to the case of no folding. The caveat of large folding, however, is that large LUT could also lower the maximum operational frequency of the design, thereby further lowering its overall performance. These effects are discussed in greater detail in Chapter ?? . Figure 7 shows a single processing element (multiplication node) with folding factor  $f$  in our implementation.

Since a number of our target applications require computation of several iterations of boolean vector multiplications, viz.  $A.v, A^2.v, \dots, A^k.v$ , given an input vector  $v$ , we provision our processing elements to support iterative matrix vector product. The number of such iterations is again provided as parameter by the user. Support for iterative multiplication would require the processing elements to perform multiplication of matrix  $A$  with input vector  $v'_k$  in the  $k$ th iteration, where  $v'_k$  is the resulting product vector after the  $(k-1)$ th iteration. Since the resulting sub-vectors  $v'_k$  are local to the processing elements in the  $k$ th iteration, this task seems trivial. However, iterative multiplication poses a unique challenge of *synchronization* of processing nodes after every iteration. This means that a processing element should only commence the  $k$ th iteration after ensuring that all other processing elements are also ready for starting  $k$ th iteration. If this is not the case and a processing element (which has already sent and received all messages of previous iteration) jumps to next iteration before others, the messages sent by this PE could be misinterpreted by other PEs to belong to previous iteration, resulting in incorrect computation of product values. One solution could be to use additional bits in messages to indicate the *id* for the iteration to which the message belongs. This would allow PEs to proceed to next iteration without waiting for others but this is expensive as the flit width is limited and this would also require buffering of incoming messages of future iterations. Other option could also be to use special messages to broadcast information to all PEs whenever a PE has completed an iteration. Each PE would maintain a list of completed PEs and would proceed to next iteration only after ensuring all others are also ready to proceed. While this seems more feasible, it also has performance penalty associated with the broadcast messages after every iteration.

In our approach, we observe that the behaviour of the network after in each iteration would be exactly identical if the PEs are exactly synchronized. Therefore, if we know the maximum number of cycles required for a PE to complete an iteration, all PEs could wait to synchronize at this cycle count after every iteration, thereby eliminating the need for additional broadcast messages. It is important to note that this count is dependent on the topology and routing policy used. We specify the count as a parameter to the processing element in our design, which could be determined by the user by simulating one iteration of the multiplication on the software. Our processing elements display the cycle counts at which the processing elements are done with one iteration (when they have sent and received all messages pertaining to that iteration), and the maximum of this count can be specified in that parameter. The source code of our processing element is included in ?? . Other parameters used by this module are specified and explained in ?? . Figure 8 shows the top level user module along with their module names in our design, consisting of several instances of processing elements connected by a CONNECT generated NoC. It must be noted that the NoC can itself be partitioned using our script, in




Fig. 8: Top module for Boolean Matrix Vector Multiplication (BMVM).

which case the top module would consist of multiple instances for NoC corresponding to different partitions, with serial wire and flow control signals between them. We have only tested our partitioned NoC for boolean vector multiplication on a single FPGA, with wires corresponding to different partitions connected internally. Since the partitioned NoC itself has been tested successfully on separate Virtex 5 boards, we believe that multiplier design should also work for multi-FPGA case. The interface to the top module is kept as simple as providing the input vector followed by a reset signal as input, and obtaining the output vector on the done signal at the output.

In our final step, we wish to encapsulate our multiplication hardware within a software library to make it usable for target applications, mostly running on software. For this purpose, we use Reusable Integration Framework for FPGA Accelerators (RIFFA 2.0) [17], which provides a simple open-source integration framework for FPGA-CPU communication over high-speed PCIe links. On the software end, RIFFA provides a simple API to open/close an FPGA connection, send/receive data from FPGA and to send reset the FPGA user module. Although we use a single channel (corresponding to network sockets and parallel user cores), RIFFA can support upto 12 channels. The data transfers in RIFFA occur through PCIe transactions, with FPGA acting as the DMA bus master. For more details, the reader can refer to [17].

We demonstrate the usage using a simple C application shown in ?? . The application resets the FPGA and sends an input vector, on which the multiplication is performed by the FPGA, before waiting to receive the output product. Since the application is written for a 32-bit processor, appropriate precautions are taken to make it compatible with the *little endianess* of the RIFFA hardware modules supporting 64-bits. We also implement the necessary glue logic for the hardware interface between RIFFA endpoint and multiplier module as shown in Figure 9. The code for this module is available in ?? .

### C. Software implementation

For comparison, we have also implemented a multi-threaded software implementation of the described boolean matrix vector



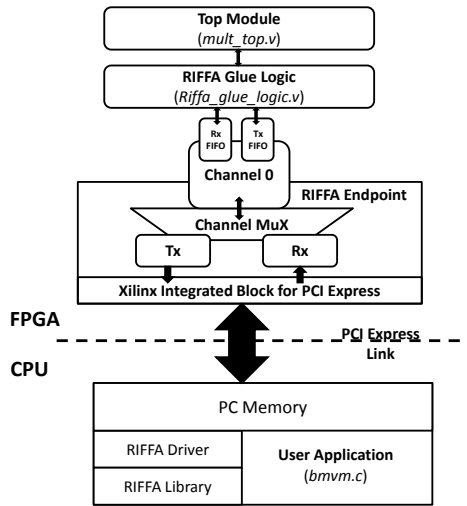


Fig. 9: BMVM framework with RIFFA.

multiplication (BMVM). The source code is available in ?? . A shared memory model, using the *pthread* library is used , in which a thread is an exact equivalent of the processing element in the FPGA. The input parameters to the application include the input matrix size, number of parallel threads, size of look-up per vector (sub-vector size to the power of 2), folding factor and the number of iterations. The threads are first initialized with look-up tables, generated using the MATLAB code as described earlier in the section, and are also provided with the input sub-vector. Since multiple parallel threads may try to update the same sub-vector at once leading to concurrency issues described earlier, mutual exclusion between parallel updates is ensured using *locks*. All the threads are also required to synchronize on a *barrier*, before proceeding to the next iteration. Our implementation, in its current form, is naive and has scope for improvement. For instance, an entire 32-bit integer is allocated for a single bit in a matrix/vector, although bit-level operations are also possible in most processors. However, even with these software level optimizations, we do not expect the application speedup by more than a few factors.

## V. CONCLUSION AND FUTURE WORK

### REFERENCES

- [1] M. K. Papamichael and J. C. Hoe, "Connect: Re-examining conventional wisdom for designing nocs in the context of fpgas," in *Proceedings of the ACM/SIGDA international symposium on Field Programmable Gate Arrays*. ACM, 2012, pp. 37–46.
- [2] I. Ouass, S. Govindarajan, V. Srinivasan, M. Kaul, and R. Vemuri, "An integrated partitioning and synthesis system for dynamically reconfigurable multi-fpga architectures," in *Parallel and Distributed Processing*. Springer, 1998, pp. 31–36.
- [3] K. Roy-Neogi and C. Sechen, "Multiple fpga partitioning with performance optimization," in *Field-Programmable Gate Arrays, 1995. FPGA'95. Proceedings of the Third International ACM Symposium on*. IEEE, 1995, pp. 146–152.
- [4] Y. Liu, P. Liu, Y. Jiang, M. Yang, K. Wu, W. Wang, and Q. Yao, "Building a multi-fpga-based emulation framework to support networks-on-chip design and verification," *International Journal of Electronics*, vol. 97, no. 10, pp. 1241–1262, 2010.
- [5] A. K. Lenstra, H. W. Lenstra Jr, M. S. Manasse, and J. M. Pollard, "The number field sieve," in *Proceedings of the twenty-second annual ACM symposium on Theory of computing*. ACM, 1990, pp. 564–572.
- [6] D. J. Bernstein, "Circuits for integer factorization: a proposal," *At the time of writing available electronically at <http://cr.yp.to/papers/nfscircuit.pdf>*, 2001.

- [7] W. Geiselmann and R. Steinwandt, "Hardware to solve sparse systems of linear equations over  $gf(2)$ ," in *Cryptographic Hardware and Embedded Systems-CHES 2003*. Springer, 2003, pp. 51–61.
- [8] S. Bajracharya, D. Misra, K. Gaj, and T. El-Ghazawi, "Reconfigurable hardware implementation of mesh routing in number field sieve factorization," in *Field-Programmable Technology, 2004. Proceedings. 2004 IEEE International Conference on*. IEEE, 2004, pp. 263–270.
- [9] Ü. V. Çatalyürek and C. Aykanat, "Decomposing irregularly sparse matrices for parallel matrix-vector multiplication," in *Parallel Algorithms for Irregularly Structured Problems*. Springer, 1996, pp. 75–86.
- [10] S. Toledo, "Improving the memory-system performance of sparse-matrix vector multiplication," *IBM Journal of research and development*, vol. 41, no. 6, pp. 711–725, 1997.
- [11] L. Barnault and D. Declercq, "Fast decoding algorithm for ldpc over  $gf(2q)$ ," in *Information Theory Workshop, 2003. Proceedings. 2003 IEEE*, March 2003, pp. 70–73.
- [12] M. D. Swanson and A. H. Tewfik, "A binary wavelet decomposition of binary images," *Image Processing, IEEE Transactions on*, vol. 5, no. 12, pp. 1637–1650, 1996.
- [13] M. J. Wiener, "Cryptanalysis of short rsa secret exponents," *Information Theory, IEEE Transactions on*, vol. 36, no. 3, pp. 553–558, 1990.
- [14] C. Anand and S. II, "Factoring of large numbers using number field sieve-the matrix step," 2007.
- [15] D. Coppersmith, "Solving homogeneous linear equations over  $(2)$  via block wiedemann algorithm," *Mathematics of Computation*, vol. 62, no. 205, pp. 333–350, 1994.
- [16] R. Williams, "Matrix-vector multiplication in sub-quadratic time:(some preprocessing required)," in *Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms*. Society for Industrial and Applied Mathematics, 2007, pp. 995–1001.
- [17] M. Jacobsen and R. Kastner, "Riffa 2.0: A reusable integration framework for fpga accelerators," in *Field Programmable Logic and Applications (FPL), 2013 23rd International Conference on*. IEEE, 2013, pp. 1–8.